

Sink Pipes API

The Sink Pipes API allows you to create, delete, schedule, and trigger Sink Pipes.

POST /v0/pipes/{pipe_id}/nodes/{node_id}/sink

Set the pipe as a Sink pipe, optionally scheduled. Required token permission: Admin.

Restrictions

- You can set only one schedule per Sink pipe.
- You can't set a Sink pipe if the pipe is already materializing. You must unlink the Materialization first.
- You can't set a Sink pipe if the pipe is already an endpoint. You must unpublish the endpoint first.

Example

```
curl \
  -X POST "https://api.tinybird.co/v0/pipes/:pipe/nodes/:node/sink" \
  -H "Authorization: Bearer <ADMIN token>" \
  -d "connection=my_connection_name" \
  -d "path=s3://bucket-name/prefix" \
  -d "region=eu-west-3" \ # Only a required parameter for IAM role
  -d "
=exported_file_template" \
  -d "format=csv" \
  -d "compression=gz" \
  -d "schedule_cron=0 */1 * * *"
```

Request parameters

KEY	TYPE	DESCRIPTION
connection	String	Name of the connection to holding the credentials to run the sink
path	String	Object store prefix into which the sink will write data

KEY	TYPE	DESCRIPTION
file_template	String	File template string. See file template for more details
format	String	Optional. Format of the exported files. Default: CSV
compression	String	Optional. Compression of the output files. Default: None
schedule_cron	String	Optional. The sink's execution schedule, in crontab format.

Successful response example

```
{
  "id": "t_529f46626c324674b3a84cd820ac2649",
  "name": "p_teste",
  "description": null,
  "endpoint": null,
  "created_at": "2024-01-18 12:57:36.503834",
  "updated_at": "2024-01-18 13:01:21.435012",
  "parent": null,
  "type": "sink",
  "last_commit": {
    "content_sha": "",
    "path": "",
    "status": "changed"
  },
  "sink_node": "t_6e8afdb8c691459b80e16541433f951b",
  "schedule": {
```



Response codes

CODE	DESCRIPTION
200	OK
404	Pipe, Node, or data connector not found, bucket doesn't exist
403	Limit reached, Query includes forbidden keywords, Pipe is already a Sink Pipe, cannot assume role
401	Invalid credentials (from connection)
400	Invalid or missing parameters, bad ARN role, invalid region name

DELETE /v0/pipes/{pipe_id}/nodes/{node_id}/sink

Removes the Sink from the pipe. This does not delete the pipe nor the node, only the sink configuration and any associated settings.

Example

```
curl \
-X DELETE "https://api.tinybird.co/v0/pipes/$1/nodes/$2/sink" \
-H "Authorization: Bearer <ADMIN token>"
```

Successful response example

```
{
  "id": "t_529f46626c324674b3a84cd820ac2649",
  "name": "p_teste",
  "description": null,
  "endpoint": null,
  "created_at": "2024-01-18 12:57:36.503834",
  "updated_at": "2024-01-19 09:27:12.069650",
  "parent": null,
  "type": "default",
  "last_commit": {
    "content_sha": "",
```

Response codes

CODE	DESCRIPTION
200	OK
404	Pipe, Node, or data connector not found
403	Limit reached, Query includes forbidden keywords, Pipe is already a Sink Pipe
400	Invalid or missing parameters, Pipe is not a Sink Pipe

POST /v0/pipes/{pipe_id}/sink

Triggers the sink pipe, creating a sink job. Allows overriding some of the sink settings for this particular execution

Example



```
curl \
-X POST "https://api.tinybird.co/v0/pipes/p_teste/sink" \
-H "Authorization: Bearer <ADMIN token>" \
-d "format=csv" \
-d "compression=gz" \
-d "file_template=export_file" \
-d <key>=<val>
```

Request parameters

KEY	TYPE	DESCRIPTION
connection	String	Name of the connection to holding the credentials to run the sink
path	String	Object store prefix into which the sink will write data
file_template	String	File template string. See file template for more details
format	String	Optional. Format of the exported files. Default: CSV
compression	String	Optional. Compression of the output files. Default: None
<key>	String	Optional. Additional variables to be injected into the file template. See file template for more details

Successful response example

```
{
  "id": "t_6e8afdb8c691459b80e16541433f951b",
  "name": "p_teste_0",
  "sql": "SELECT * FROM teste",
  "description": null,
  "materialized": null,
  "cluster": null,
  "tags": {},
  "created_at": "2024-01-18 12:57:36.503843",
  "updated_at": "2024-01-19 09:27:12.069649",
  "version": 0,
  "project": null,
  "result": null,
  "ignore_sql_errors": false,
```



Response codes

CODE	DESCRIPTION
200	OK
404	Pipe, Node, or data connector not found
403	Limit reached, Query includes forbidden keywords, Pipe is already a Sink Pipe
400	Invalid or missing parameters, Pipe is not a Sink Pipe

GET /v0/integrations/s3/policies/trust-policy

Retrieves the trust policy to be attached to the IAM role that will be used for the connection. External IDs are different for each Workspace, but shared between Branches of the same Workspace to avoid having to change the trust policy for each Branch.

Example

```
curl \
-X GET "https://$TB_HOST/v0/integrations/s3/policies/trust-policy" \
-H "Authorization: Bearer <ADMIN token>"
```

Successful response example

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "sts:AssumeRole",
      "Principal": {
        "AWS": "arn:aws:iam::123456789:root"
      },
      "Condition": {
```

Response codes

CODE	DESCRIPTION
200	OK
404	S3 integration not supported in your region

GET /v0/integrations/s3/policies/access-policy

Retrieves the trust policy to be attached to the IAM Role that will be used for the connection. External IDs are different for each workspace, but shared between branches of the same workspace to avoid having to change the trust policy for each branch.

Example

```
curl \
  -X GET "https://$TB_HOST/v0/integrations/s3/policies/access-policy?bucket=test"
  -H "Authorization: Bearer <ADMIN token>"
```

Request parameters

KEY	TYPE	DESCRIPTION
bucket	Optional[String]	Bucket to use for rendering the template. If not provided the ' <code><bucket></code> ' placeholder is used

Successful response example

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "s3:GetBucketLocation",
        "s3:ListBucket"
      ],
      "Resource": "arn:aws:s3:::<bucket>"
    },
    {
      "Effect": "Allow",
      "Action": [
```

```
"s3:GetObject",  
"s3:PutObject",
```

Response codes

CODE	DESCRIPTION
200	OK

GET /v0/datasources-bigquery-credentials

Retrieves the workspace’s GCP service account to be authorized to write to the destination bucket.

Example

```
curl \  
-X POST "${TINYBIRD_HOST}/v0/connectors" \  
-H "Authorization: Bearer <ADMIN TOKEN>" \  
-d "service=gcs_service_account" \  
-d "name=<name>"
```



Request parameters

None

Successful response example

```
{  
  "account": "cdk-E-d83f6d01-b5c1-40-43439d@development-353413.iam.gserviceacc  
}
```



Response codes

CODE	DESCRIPTION
200	OK
503	Feature not enabled in your region

Company

- Product
- Pricing
- ROI Calculator
- About Us
- Shop
- Careers
- Request a demo

Integrations

- Amazon S3
- Kafka Data Streams
- Google Cloud Storage
- Google BigQuery
- Snowflake
- Confluent

Resources

- Docs
- Blog
- Community
- Live Coding
- Customer Stories
- RSS Feed

Use cases

- In-Product Analytics
- Operational Intelligence
- Realtime Personalization
- Anomaly Detection & Alerts
- Usage Based Pricing
- Sports Betting/Gaming
- Smart Inventory Management
- Serverless ClickHouse

Copyright © 2024 Tinybird. All rights reserved
[Terms & conditions](#) [Cookies](#) [Security](#)

Spain
Calle del Dr. Fourquet, 27
28012 Madrid

USA
41 East 11th Street 11th floor
New York, NY 10003